MLOps Assignment 1

Deep Learning Experiments on MNIST and FashionMNIST

Roll Number: B22CSXXX
Name: FirstName LastName

January 24, 2026

Google Colab Link

Colab Notebook: `https://colab.research.google.com/YOUR_COLAB_LINK_HERE`

GitHub Repository: `https://github.com/YOUR_USERNAME/MLOps-Name-RollNumber`

Contents

1 Introduction

This report presents the results of deep learning experiments conducted on the MNIST and Fashion-MNIST datasets using ResNet architectures. The experiments cover:

- **Q1(a):** Training ResNet-18 and ResNet-50 models with various hyperparameters
- **Q1(b):** SVM classifier baseline with polynomial and RBF kernels
- **Q2:** CPU vs GPU performance comparison

1.1 Experimental Setup

- **Data Split:** 70% Train - 10% Validation - 20% Test
- **Models:** ResNet-18, ResNet-50 (pretrained=False)
- **USE_AMP:** True
- **pin_memory:** True, False (varied)
- **Epochs:** 2, 5 (varied)

2 Q1(a): ResNet Classification Results

2.1 MNIST Dataset

Table 1: MNIST Test Classification Accuracy (%)

Batch Size	Optimizer	Learning Rate	ResNet-18	ResNet-50
16	SGD	0.001	98.95	98.55
16	SGD	0.0001	98.48	97.91
16	Adam	0.001	**99.01**	98.31
16	Adam	0.0001	98.76	97.74
32	SGD	0.001	98.82	**98.70**
32 [Optional]	SGD	0.0001	97.91	96.87
32 [Optional]	Adam	0.001	98.78	98.54
32	Adam	0.0001	98.66	97.02

2.2 FashionMNIST Dataset

Table 2: FashionMNIST Test Classification Accuracy (%)

Batch Size	Optimizer	Learning Rate	ResNet-18	ResNet-50
16	SGD	0.001	90.06	89.02
16	SGD	0.0001	88.66	85.60
16	Adam	0.001	90.57	87.42
16	Adam	0.0001	89.35	87.91
32	SGD	0.001	89.93	89.29
32 [Optional]	SGD	0.0001	87.86	83.60
32 [Optional]	Adam	0.001	**90.76**	**89.40**
32	Adam	0.0001	89.69	86.91

2.3 Additional Hyperparameter Experiments

Experiments were also conducted varying `pin_memory` (True/False) and epochs (2, 5).

Table 3: Effect of Epochs and Pin Memory on Performance (MNIST, ResNet-18)

BS	Opt	LR	Pin Mem	Epochs	Accuracy (%)
16	Adam	0.001	False	2	98.36
16	Adam	0.001	False	5	99.01
16	Adam	0.001	True	2	98.32
16	Adam	0.001	True	5	98.52
16	SGD	0.001	False	2	98.46
16	SGD	0.001	False	5	98.89
16	SGD	0.001	True	2	98.59
16	SGD	0.001	True	5	98.95

2.4 Analysis - Q1(a)

Key Observations:

1. **ResNet-18 vs ResNet-50:** Despite having fewer parameters (11.18M vs 23.52M), ResNet-18 consistently outperforms ResNet-50 on both datasets. This indicates that for relatively simple image classification tasks like MNIST and FashionMNIST, a lighter model architecture is sufficient and may generalize better due to reduced overfitting.

2. **Learning Rate Impact:** A learning rate of 0.001 consistently outperforms 0.0001 across all configurations. The higher learning rate enables faster convergence and better final accuracy within the limited training epochs.

3. **Optimizer Comparison:**
   - Adam achieves the best peak accuracy (99.01% on MNIST)
   - SGD with lr=0.001 is competitive and sometimes more stable
   - Adam with lr=0.0001 performs well due to adaptive learning rates

4. **Batch Size Effect:** The difference between batch sizes 16 and 32 is minimal ($< 1\%$), suggesting both are suitable for these datasets.

5. **Epochs:** Increasing from 2 to 5 epochs improves accuracy by 0.3-0.7%, indicating the models benefit from additional training.

6. **Dataset Complexity:** MNIST achieves much higher accuracy ($\sim$99%) compared to Fashion-MNIST ($\sim$91%), reflecting the greater visual complexity of fashion items versus handwritten digits.

Best Configurations:

- MNIST: ResNet-18, Adam, lr=0.001, bs=16, 5 epochs $\rightarrow$ **99.01%**
- FashionMNIST: ResNet-18, Adam, lr=0.001, bs=32, 5 epochs $\rightarrow$ **90.76%**

3 Q1(b): SVM Classification Results

SVM classifiers were trained as baselines using polynomial and RBF kernels with varying hyperparameters.

3.1 MNIST SVM Results

Table 4: SVM Results on MNIST Dataset

Kernel	C	Gamma	Test Acc (%)	F1 Score	Train Time (ms)
poly	0.1	scale	54.75	0.579	39,593
poly	1.0	scale	91.50	0.919	22,815
poly	10.0	scale	95.50	0.955	16,030
poly	0.1	auto	38.30	0.415	41,302
poly	1.0	auto	87.95	0.888	25,879
poly	10.0	auto	95.35	0.954	16,236
rbf	0.1	scale	89.35	0.895	20,014
rbf	1.0	scale	94.85	0.948	11,362
rbf	10.0	scale	**95.60**	**0.956**	10,730
rbf	0.1	auto	89.85	0.899	19,744
rbf	1.0	auto	95.20	0.952	10,718
rbf	10.0	auto	95.55	0.955	9,776

3.2 FashionMNIST SVM Results

Table 5: SVM Results on FashionMNIST Dataset

Kernel	C	Gamma	Test Acc (%)	F1 Score	Train Time (ms)
poly	0.1	scale	69.05	0.716	21,222
poly	1.0	scale	82.15	0.827	11,343
poly	10.0	scale	85.15	0.852	9,358
poly	0.1	auto	69.05	0.716	21,333
poly	1.0	auto	82.15	0.827	11,508
poly	10.0	auto	85.15	0.852	9,229
rbf	0.1	scale	78.65	0.783	14,453
rbf	1.0	scale	85.00	0.849	9,353
rbf	10.0	scale	**86.30**	**0.862**	9,091
rbf	0.1	auto	78.65	0.783	14,830
rbf	1.0	auto	85.00	0.849	9,144
rbf	10.0	auto	**86.30**	**0.862**	9,033

3.3 SVM vs Deep Learning Comparison

Table 6: Best Results: SVM vs ResNet Comparison

Dataset	Method	Configuration	Accuracy (%)	F1 Score
MNIST	SVM	rbf, C=10.0, gamma=scale	95.60	0.956
MNIST	ResNet-18	Adam, lr=0.001, bs=16	**99.01**	**0.990**
FashionMNIST	SVM	rbf, C=10.0, gamma=scale	86.30	0.862
FashionMNIST	ResNet-18	Adam, lr=0.001, bs=32	**90.76**	**0.907**

3.4 Analysis - Q1(b)

Key Observations:

1. **Kernel Performance:** RBF kernel consistently outperforms polynomial kernel, especially for FashionMNIST. RBF is better suited for image classification due to its ability to model complex decision boundaries.

2. **Regularization (C):** Higher C values (C=10.0) achieve better accuracy but may risk overfitting. C=10.0 provides the best trade-off for both datasets.

3. **Gamma Setting:** "scale" generally performs slightly better than "auto" for polynomial kernels, while the difference is minimal for RBF.

4. **Training Time:** SVM training time decreases with higher C values due to fewer support vectors. Best SVM takes ∼10 seconds compared to minutes for deep learning.

5. **Deep Learning Superiority:** ResNet models significantly outperform SVM:
   - MNIST: +3.41% improvement (99.01% vs 95.60%)
   - FashionMNIST: +4.46% improvement (90.76% vs 86.30%)

4 Q2: CPU vs GPU Performance Comparison

Performance comparison between CPU and GPU (CUDA) training on FashionMNIST dataset.

4.1 Experimental Configuration

- **Dataset:** FashionMNIST
- **Batch Size:** 32
- **Learning Rate:** 0.001
- **Epochs:** 5

4.2 Performance Results

Table 7: CPU vs GPU Performance Comparison on FashionMNIST

Compute	BS	Opt	LR	Test Acc (%)		Train Time (ms)		FLOPs	
				ResNet-18	ResNet-50	ResNet-18	ResNet-50	ResNet-18	ResNet-50
CPU	32	SGD	0.001	88.54	84.84	341,621	1,064,108	33.18M	78.76M
CPU	32	Adam	0.001	88.23	87.11	470,803	1,281,028	33.18M	78.76M
CUDA	32	SGD	0.001	87.97	84.18	62,967	111,651	33.18M	78.76M
CUDA	32	Adam	0.001	86.14	86.96	69,794	125,693	33.18M	78.76M

4.3 GPU Speedup Analysis

Table 8: GPU Speedup Over CPU

Model	Optimizer	CPU Time (ms)	GPU Time (ms)	Speedup
ResNet-18	SGD	341,621	62,967	**5.43x**
ResNet-18	Adam	470,803	69,794	**6.75x**
ResNet-50	SGD	1,064,108	111,651	**9.53x**
ResNet-50	Adam	1,281,028	125,693	**10.19x**

4.4 Model Complexity

Table 9: Model Complexity Comparison

Model	Parameters	FLOPs
ResNet-18	11.18M	33.18M
ResNet-50	23.52M	78.76M

4.5 Analysis - Q2

Training Time Analysis:

1. **GPU Acceleration:** GPU provides significant speedup over CPU:
   - ResNet-18: 5.43x - 6.75x faster on GPU
   - ResNet-50: 9.53x - 10.19x faster on GPU
2. **Model Complexity Impact:** Larger models benefit more from GPU acceleration. ResNet-50 achieves nearly 10x speedup compared to 5-7x for ResNet-18.
3. **Optimizer Overhead:** Adam takes longer than SGD on both CPU and GPU due to maintaining additional moment estimates. On CPU, Adam is ∼38% slower than SGD for ResNet-18.

FLOPs Analysis:

1. ResNet-50 requires 2.37x more FLOPs than ResNet-18 (78.76M vs 33.18M)
2. This correlates with the training time difference: ResNet-50 takes approximately 2x longer to train
3. FLOPs remain constant regardless of CPU/GPU—only execution time differs

Classification Accuracy Analysis:

1. Slight accuracy variations between CPU and GPU runs are due to numerical precision differences and random initialization
2. CPU training achieved marginally higher accuracy in some cases, likely due to different floating-point handling
3. Both achieve similar final performance, confirming that GPU acceleration doesn't sacrifice model quality

Practical Implications:

- For rapid experimentation, GPU training is essential—reduces training from minutes to seconds
- For deployment on resource-constrained devices, ResNet-18 offers the best accuracy-efficiency trade-off
- Adam optimizer provides faster convergence but at higher computational cost

5 Conclusions

5.1 Summary of Results

Table 10: Best Configurations Summary

Task	Best Configuration	Accuracy
MNIST (Deep Learning)	ResNet-18, Adam, lr=0.001, bs=16	99.01%
FashionMNIST (Deep Learning)	ResNet-18, Adam, lr=0.001, bs=32	90.76%
MNIST (SVM)	RBF, C=10.0, gamma=scale	95.60%
FashionMNIST (SVM)	RBF, C=10.0, gamma=scale	86.30%

5.2 Key Takeaways

1. **Model Selection:** ResNet-18 is the optimal choice for MNIST/FashionMNIST—it outperforms the larger ResNet-50 while being faster to train.

2. **Hyperparameter Tuning:** Learning rate of 0.001 with Adam optimizer provides the best results. Batch size has minimal impact.
3. **Deep Learning vs Traditional ML:** ResNet models significantly outperform SVM classifiers (+3-5% accuracy), justifying the additional computational cost.
4. **GPU Utilization:** GPU training provides 5-10x speedup, making it essential for efficient deep learning experimentation.
5. **All experiments achieved >80% classification accuracy**, meeting the assignment requirements.